DEPARTMENT OF INFORMATICS

TECHNICAL UNIVERSITY OF MUNICH

Bachelor's Thesis in Informatics

Formalisation of a Congruence Closure Algorithm in Isabelle/HOL

Rebecca Ghidini

DEPARTMENT OF INFORMATICS

TECHNICAL UNIVERSITY OF MUNICH

Bachelor's Thesis in Informatics

Formalisation of a Congruence Closure Algorithm in Isabelle/HOL

Formalisierung eines Kongruenzhüllen-Algorithmus in Isabelle/HOL

Author:	Rebecca Ghidini
Supervisor:	Prof. Dr. Tobias Nipkow
Advisor:	Lukas Stevens
Submission Date:	15.09.2022

I confirm that this bachelor's thesis in informatics is my own work and I have documented all sources and material used.

Munich, 15.09.2022   Rebecca Ghidini

Acknowledgments

Thanks to Timmm and Manon.

Abstract

Interactive theorem provers are useful for the formalization and proving of formulas. Congruence closure is already used by most theorem provers for solving decision procedures. However, the implementations of congruence closure are usually not formally proven for completeness.

This thesis describes the implementation and the proofs for termination and correctness of an explain operation for union-find and of a congruence closure algorithm, whose implementation is partly based on the union-find algorithm. It also contains an implementation of the explain operation for congruence closure. The proofs are formalized in the interactive theorem prover Isabelle/HOL. The implementation is based on the paper by Nieuwenhuis and Oliveras [1] and on the formalization of the union-find data structure in Isabelle/HOL by Lammich [2]. The implementation can be used in the future in order to develop an automatic proof strategy in Isabelle, for automatic proofs of theorems.

Contents

1 Introduction

The equivalence closure of a relation is the smallest superset of the relation which is reflexive, symmetric and transitive. For example, $a = c$ is in the equivalence closure of $a = b$ and $b = c$. Algorithms for maintaining the equivalence closure have been known for a long time [3], the most widely used being the union-find algorithm due to its simplicity and almost constant runtime [4].

The congruence closure of a set of equations is similar, but it considers functions in addition to constants. Further to reflexivity, symmetry and transitivity, it satisfies monotonicity, i.e. $f(x_1, ..., x_n) = f(y_1, ..., y_n)$ if $x_i = y_i$ for all i between 1 and n. [1] For example, the equation $c = e$ belongs to the congruence closure of $f(a, b) = c$, $f(d, b) = e$ and $a = d$. Several approaches to solve this problem have been described, with different runtimes and applications, as shown by [5, 6, 7, 1].

These algorithms are used in decision procedures such as satisfiability modulo theories (SMT) solvers. [8] In these settings, it is required to generate an explanation by finding the set of input equations which caused the congruence. These can then be used to generate a certificate of the congruence, which can be verified by a program. Nieuwenhuis and Oliveras have presented an efficient version of the congruence closure algorithm and two versions of the union-find algorithm, each with their own explain operation, in a conference paper [1], which was later extended, see [9].

The descriptions of the algorithms contain informal proofs, but the algorithms can be verified by an interactive theorem prover in order to strengthen our confidence in their correctness. In this thesis we will implement the algorithms of the paper [1] and prove their correctness in the theorem prover Isabelle/HOL. The description of the congruence closure algorithms contains an informal proof, but to my knowledge, this thesis presents the first proof of the union-find explain algorithm and of the congruence closure algorithm in an automatic theorem prover. Our implementation is based on the union-find formalization by Lammich [2] in Isabelle/HOL. Given that the focus of this paper is on the verification of the algorithms, a few optimizations are left out of the implementation, such as path compression for union-find.

1.1 Outline

This thesis is organized as follows: Chapter 2 discusses some related work and gives a brief overview of the notation used by Isabelle.

In Chapter 3 the union-find implementation by Lammich [2] is described, and the explain operation is presented together with its correctness and termination proofs.

Chapter 4 looks at the congruence closure implementation and its correctness and termination proof. It also describes the explain operation for congruence closure, and discusses a possible proof outline for its correctness. For reasons of time, the actual proof has not been finished yet.

The last chapter summarizes the results and gives an outlook on the proofs which are open for future work.

The Isabelle code of this thesis is available on GitHub[1].

---

[1]https://github.com/reb-ddm/congruence-closure-isabelle

2 Preliminaries

2.1 Isabelle/HOL

Isabelle[10] is an automated theorem prover, which can be used among other things for the verification of algorithms. It provides different types of logic, the most used one being Higher-Order Logic (HOL), which is the one used in this thesis. For reasons of self-containedness of this thesis, this chapter introduces the notation used by Isabelle/HOL.

Lists

The syntax for the empty list in Isabelle is `[]`, and longer lists are constructed with the infix # operator, with which one element is appended to the front of the list. To concatenate two lists, the @ operator is used. `set` is a function which converts a list to a set. Lists are indexed with the ! operator, and the syntax for updating a list l at index i is the following: `l[i := new_value]`.

Functions

In Isabelle, the termination of functions must be proven. In the case of simple functions, Isabelle can prove it automatically. This is done for functions which are declared with the **fun** keyword. For example the declaration **fun** `f :: 'a` $\Rightarrow$ `'b` describes a function f with a parameter of the type $'a$ and it returns a value of the type $'b$. Afterwards, the recursive equations of the function are defined. Isabelle automatically defines induction rules for each function.

Partial functions can also be defined in Isabelle, with the **function** keyword. The definition looks like this:

```
function (domintros) g
  where "..."
  by pat_completeness auto
```

Isabelle automatically defines a predicate `g_dom` where `g_dom (a)` means that the function `g` terminates with the parameter a.

The option `domintros` provides inductive introduction rules for `g_dom`, based on the defining equations of g.

After the function definition, it needs to be proven that the patterns used in the definition are complete and compatible. In our case, the method `pat_completeness` `auto` always automatically proves this goal.

Pertial simplification rules and a partial induction rule are also automatically defined by Isabelle, they can only be applied if we assume or prove that the function terminates with the given parameters.

For a more detailed description of functions in Isabelle, see [11].

Records

Records are similar to tuples, where each component has a name. For example, for the implementation of the union-find explain operations, we need three lists, which are grouped together in a record. The syntax is `ufe = ⦇uf_list = l, unions = u, au = a⦈`, and to select for example the first component, we can write `uf_list ufe`. The meaning of the three lists will be described in Section 3.3.

For more information on records, see [10], chapter 8.3.

equivalence closure

The theory "Partial_Equivalence_Relation"[12] defines the symmetric closure `symcl` of a relation and the reflexiv-transitive closure is already part of the Isabelle/HOL distribution, and its syntax is `R*`. The two definitions can be combined to have an abstraction of the equivalence closure. For example the equivalence closure of the relation R is `(symcl R)*`.

Datatypes

New datatypes can be defined with the **`datatype`** keyword. New datatypes consist of constructors and existing types. A concrete syntax for the new datatypes can be defined in brackets.

For example we define a new datatype for the two types of input equations used in the congruence closure algorithm. Equations of the type $a = b$ will be written as `a ≈ b`, and equations of the type $F(a, b) = c$ are written as `F a b ≈ c`.

```
datatype equation = Constants nat nat ("_ ≈ _")
  | Function nat nat nat ("F _ _ ≈ _")
```

In this thesis, we will use the notation `a ≈ b` and `F a b ≈ c` in Isabelle listings, and $a = b$ and $F(a, b) = c$ outside of the listings.

option

The type `option` models optional values. The value of a variable with type `'a option` is either `None` or `Some a` where *a* is a value with type $'a$. The function `the` applied to `Some a` returns `a`, and it returns `undefined` if the parameter is `None`.

If `'a` is an ordered type, the order is extended to `'a option`, where `None` ≤ `x` for all `x`, and `Some x` ≤ `Some y` iff `x` ≤ `y`. This is defined in the Theory "Option_ord" of the HOL library, which is included with the standard Isabelle/HOL distribution.

2.1.1 Related work

Efficient union-find algorithms have been known for a long time, see [3, 4]. Given its importance as an algorithm, it was already formalized and verified in some of the most important theorem provers, such as Isabelle and Coq [13]. The code in this thesis uses the union-find formalization in Isabelle by Lammich, which was first published in a journal [2] and later presented at a conference [14]. It includes the functions for *union* and *find*, as well as an invariant which charachterizes the validity of the union-find data structure. It will be described in more detail in Section 3.2.

Based on the union-find implementation, efficient congruence closure algorithms have been developed by Shostak [5], Nelson and Oppen [6] and Downey at al. [7]. Nieuwenhuis and Oliveras [1] extend the algorithm by an *explain* operation, which is necessary in the context of decision procedures, for example for theorem provers.

Congruence closure is implemented in most automatic theorem provers. Pierre Corbineau implemented a `congruence` tactic for the theorem prover Coq [15, 16], based on the algorithm of Downey at al. [7], and with an *explain* operation which is similar to the union-find explain presented in this thesis, and not as efficient as the *explain* operation introduced by Nieuwenhuis [1].

Lean also has a congruence-closure based decision procedure [17], which is additionally able to handle dependent types.

Isabelle/HOL includes a tool called `sledgehammer` [18], which uses external SMT solvers, e.g. Z3 [8] and CVC4 [19], whose implementation is based on congruence closure. However, there is no built-in congruence closure proof method for Isabelle yet. The verified algorithm of this paper can be used in the future in order to build such a proof method.

3 Union-Find with Explain Operation

3.1 Union-Find Algorithm

The union-find data structure is used to keep track of the equivalence closure of a set of equations between constants, by partitioning the initial set of n constants into equivalence classes. It is initialized with a partition where each element is in their own equivalence class. There is a `union` operation which merges two equivalence classes. The `find` operation returns a representative of the equivalence class. Two elements are in the same representative class, if and only if they have the same representative.

The equivalence classes are modeled with a forest, which is a graph where each connected component is a tree. Initially, the graph contains n vertices and no edges, then each `union` adds a directed edge. The connected components of the graph represent the equivalence classes. Each tree in the forest has a root, which is also the representative of the equivalence class, and each edge in the tree is directed towards the root. In order to keep this invariant, at each `union a b`, the new edge is added between the representative of a (which will be denoted as $rep\_of(l, a)$, where l is the union-find forest) and $rep\_of(l, b)$.

The union-find forest is modeled by a list l of length n, where at each index i the list contains the parent of the element i in the forest. If i does not have a parent, i.e. it is a root, then the list contains the element i itself at the index i, e.g. $l!i = i$.

The original union-find algorithm [4] contains two tions: path compression in the `find` method, and choosing the representative of the bigger class to be the new representative of the merged class in `union`. These optimizations are irrelevant for the correctness of the algorithm, therefore we leave them out of this implementation in order to simplify the proofs.

3.2 Union-Find in Isabelle

The union-find algorithm was already formalized in Isabelle by Lammich [2], and the code can be found in the "Archive of Formal Proofs" (AFP) under "A Separation Logic Framework for Imperative HOL", in the theory "Union_Find"[20, 21]. The following is a brief description of the implementation.

The function `rep_of` finds the representative of an element in the forest. It is analogous to the `find` operation, except that it does not do path compression.

```
function (domintros) rep_of
  where "rep_of l i = (if l!i = i then i else rep_of l (l!i))"
  by pat_completeness auto
```

The domain of `rep_of` is used to define the following invariant for valid union-find lists. This invariant states that the `rep_of` function terminates on all valid indexes of the list, which is equivalent to saying that the union-find forest does not contain any cycles.

```
definition
  "ufa_invar l ≡ ∀i<length l. rep_of_dom (l,i) ∧ l!i<length l"
```

The `union` operation simply adds an edge between the representatives of the two elements.

```
abbreviation "ufa_union l x y ≡ l[rep_of l x := rep_of l y]"
```

The theory contains several lemmas, including the lemma which states that the invariant `ufa_invar` holds for the initial union-find forest without edges, and that it is preserved by the `ufa_union` operation.

3.3 Union-Find Data Structure

The section below describes the implementation of a union-find data structure, which was modified in order to support an `explain` operation, as described in [1].

The data structure for the Union, Find and Explain operations consists of the following three lists:

- `uf_list`: This is the usual union-find list, which contains the parent node of each element in the forest data structure. It is the one described in Section 3.1.

- `unions`: This list simply contains all the pairs of input elements in chronological order.

- au: This is the *associated unions* list, it contains for each edge in the union-find forest a label with the union that corresponds to this edge. Similarly to the `uf_list`, it is indexed by the element, and for each element e which has a parent in the `uf_list`, au contains the input equation which caused the creation of this edge between e and its parent. The equations are represented as indexes in the `unions` list. The type of the entries is `nat option`, so that for elements without a parent, the au entry is `None`.

Example 1. For a union-find algorithm with 4 variables, the initial empty union find looks as follows:

⦇uf_list = [0, 1, 2, 3], unions = [], au = [None, None, None, None]⦈

Each element is its own parent in the uf_list, which means that it is a root, the unions list is empty because no unions were made yet, and there are no edges in the tree, therefore there are no labels in au.

In order to reason about paths in the union-find forest, we define the following path predicate.

```
inductive path :: "nat list ⇒ nat ⇒ nat list ⇒ nat ⇒ bool" where
  single: "n < length l ⟹ path l n [n] n" |
  step: "r < length l ⟹ l ! u = r ⟹ l ! u ≠ u ⟹ path l u p v ⟹ path l r
      (r # p) v"
```

path l r p v defines a path from r to v, where r is an ancestor of v, which means that it is closer to the root, and p is a list which contains all the nodes visited on the path from r to v. This definition proved to be very useful for many proofs, as will become clearer later in this thesis.

The theory "Path" contains many lemmas about paths, including lemmas about concatenation of adjacent paths, and splitting of one path into two subpaths, and that the length of a path is at least 1, as well as others, many of which could be proven by rule induction on path. The most interesting and useful lemma was about the uniqueness of paths between two nodes:

```
theorem path_unique: "ufa_invar l ⟹ path l u p1 v ⟹ path l u p2 v ⟹ p1 =
    p2"
```

Proof. The lemma is proven by induction on the length of $p1$.

For the base case we assume that the length of $p1$ is 1. There is only one node in the path, therefore $v = u$. Then I proved a lemma which showed that if the ufa_invar holds, each path from v to v has length 1, or, in other words, there are no cycles in the graph. For this I showed that if there was a cycle, the function rep_of would not terminate, because there would be an infinite loop.

For the induction step, we assume that the length of $p1$ is greater than 1. Therefore, we can remove the last node from $p1$ and the last node from $p2$ to get two paths $p1'$ and $p2'$ from u to the parent of v, where $p1'$ is shorter that $p1$, and we can apply the induction hypothesis, which tells us that $p1' = p2'$. Adding the node v to those two paths gives us back the original paths $p1$ and $p2$, therefore we conclude that $p1 = p2$. □

3.4 Implementation

3.4.1 Union

The `ufa_union` operation needs to be extended in order to appropiately update the other two lists:

The function `ufe_union` only modifies the data structure if the parameters are not already in the same equivalence class. The `uf_list` is modified with `ufa_union`. The current union (x, y) is appended to the end of the `unions` list. `au` is updated such that the new edge between `rep_of l x` and `rep_of l y` is labeled with the last index of `unions`, which contains the current pair of elements (x, y).

```
fun ufe_union :: "ufe_data_structure ⇒ nat ⇒ nat ⇒ ufe_data_structure"
  where
    "ufe_union ⦇uf_list = l, unions = u, au = a⦈ x y = (
if (rep_of l x ≠ rep_of l y) then
    ⦇uf_list = ufa_union l x y,
     unions = u @ [(x,y)],
     au = a[rep_of l x := Some (length u)]⦈
else ⦇uf_list = l, unions = u, au = a⦈)"
```

Example 2. After a union of 0 and 1, the data structure from Example 1 looks as follows:

```
⦇uf_list = [1, 1, 2, 3], unions = [(0, 1)], au = [Some 0, None, None, None]⦈
```

There is an edge between 1 and 0, labeled with the union at index 0, which is $(0, 1)$.

Next, we define a function which takes a list of unions as parameter and simply applies each of those unions to the data structure. This will be needed for the invariant and the correctness proof in the next sections.

```
fun apply_unions :: "(nat * nat) list ⇒ ufe_data_structure ⇒
    ufe_data_structure"
  where
    "apply_unions [] p = p" |
    "apply_unions ((x, y) # u) p = apply_unions u (ufe_union p x y)"
```

3.4.2 Helper Functions for Explain

We implement the explain function following the description of the first version of the union-find algorithm in the paper[1].

The explain function takes as parameter two elements x and y and calculates a subset of the input unions which explain why the two given variables are in the same equivalence class. If we consider the graph which has as nodes the elements and as

edges the input unions, then the output of explain would be all the unions on the path from x to y. However, the union-find forest in our data structure does not have as edges the unions, but only edges between representatives of the elements of the unions.

From this graph, we can calculate the desired output in the following way: first add the last union (a, b) made between the equivalence class of x and the one of y, then recursively call the explain operation with the new parameters (x, a) and (b, y) (or (x, b) and (a, y), depending on which branch a and b are on). The last union is the edge with the highest label in the union-find forest.

(a, b) is calculated by finding the lowest common ancestor *lca* of x and y, and then finding the newest union on the path from x to *lca* and from y to *lca*.

This section describes the helper functions needed for the implementation of `explain`, which calculate the lowest common ancestor, using the function `path_to_root`, and the newest union on a path.

path_to_root

The function `path_to_root l x` computes the path from the root of x to the node x in the union-find forest l. It simply starts at x and continues to add the parent of the current node to the front of the path, until it reaches the root.

```
function path_to_root :: "nat list ⇒ nat ⇒ nat list"
  where
    "path_to_root l x = (if l ! x = x then [x] else path_to_root l (l ! x) @ [
        x])"
  by pat_completeness auto
```

It was easy to show that it has the same domain as the `rep_of` function, as it has the same recursive calls.

```
lemma path_to_root_domain: "rep_of_dom (l, i) ⟷ path_to_root_dom (l, i)"
```

The correctness of the function follows easily by induction.

```
theorem path_to_root_correct:
  assumes "ufa_invar l"
  shows "path l (rep_of l x) (path_to_root l x) x"
```

lowest_common_ancestor

The function `lowest_common_ancestor l x y` finds the lowest common ancestor of x and y in the union-find forest l.

Definition. A *common ancestor* of two nodes x and y is a node which has a path to x and a path to y. The *lowest common ancestor* of two nodes x and y is the common ancestor which is farthest away from the root.

The function will only be used for two nodes which have the same root, otherwise there is no common ancestor. It first computes the paths from x and y to their root, and then returns the last element which the two paths have in common. For this it uses the function `longest_common_prefix` from "HOL-Library.Sublist", which is included in the standard Isabelle distribution.

```
fun lowest_common_ancestor :: "nat list ⇒ nat ⇒ nat ⇒ nat"
  where
    "lowest_common_ancestor l x y =
last (longest_common_prefix (path_to_root l x) (path_to_root l y))"
```

Regarding the correctness proof, there were two aspects to prove: the most useful result is that `lowest_common_ancestor l x y` is a common ancestor of x and y. The second aspect stated that any other common ancestor of x and y has a shorter distance from the root. The proof assumes that that x and y have the same root.

```
abbreviation "common_ancestor l x y ca ≡
(∃ p . path l ca p x) ∧
(∃ p . path l ca p y)"

abbreviation "Lowest_common_ancestor l x y ca ≡
(common_ancestor l x y ca ∧
(∀r ca2 p3 p4. (path l r p3 ca ∧ path l r p4 ca2 ∧ common_ancestor l x y ca2
⟶ length p3 ≥ length p4)))"

theorem lowest_common_ancestor_correct:
  assumes "ufa_invar l"
    and "x < length l"
    and "y < length l"
    and "rep_of l x = rep_of l y"
  shows "Lowest_common_ancestor l x y (lowest_common_ancestor l x y)"
```

Proof. Let $lca =$ `lowest_common_ancestor l x y`. We previously proved that `path_to_root` computes a path p_x from the root to x and a path p_y from the root to y. Evidently, lca lies on both paths, because it is part of their common prefix. Splitting the path p_x, we get a path from the root to lca and one from lca to x, and the same for y. This shows that lca is a common ancestor.

To prove that it is the *lowest* common ancestor, we can prove it by contradiction. We assume that there is a common ancestor lca_2 with a longer path from the root than lca, then we can show that there is a path from the root to x passing through lca_2, and the same for y. Because of the uniqueness of paths, these paths are equal to `path_to_root l x` and `path_to_root l y`, respectively. That means, that there is a prefix of `path_to_root l x` and `path_to_root l y` which is longer than the one calculated by the function `longest_common_prefix`. The theory "Sublist" contains a correctness proof for `longest_common_prefix`, which we can use to show the contradiction. □

find_newest_on_path

The function `find_newest_on_path` finds the newest edge on the path from y to x. It is assumed that y is an ancestor of x. The function simply checks all the elements on the path from y to x and returns the one with the largest index in a, which is the *associated unions* list.

```
function (domintros) find_newest_on_path :: "nat list ⇒ nat option list ⇒
    nat ⇒ nat ⇒ nat option"
  where
    "find_newest_on_path l a x y =
  (if x = y then None
    else max (a ! x) (find_newest_on_path l a (l ! x) y))"
  by pat_completeness auto
```

If there is a path p from y to x, it is easily shown by induction that the function terminates.

```
lemma find_newest_on_path_domain:
  assumes "ufa_invar l"
    and "path l y p x"
  shows "find_newest_on_path_dom (l, a, x, y)"
```

Note that some additional assumptions of the type `"x < length l"` are in the original formulation of this lemma. The assumption that all the variables are in bounds is present in all the lemmas about union-find and congruence closure, but they will not be mentioned in the thesis for reasons of conciseness. For the exact formulation of the lemmas, see the Isabelle code.

For the correctness proof we define an abstract definition of the newest element on the path: `Newest_on_path` is the maximal value in the associated unions list for indexes in p.

```
abbreviation "Newest_on_path l a x y newest ≡
∃ p . path l y p x ∧ newest = (MAX i ∈ set [1..<length p]. a ! (p ! i))"
```

Then it can easily be shown by computation induction on `find_newest_on_path` that our function is correct.

```
theorem find_newest_on_path_correct:
  assumes path: "path l y p x"
    and invar: "ufa_invar l"
    and xy: "x ≠ y"
  shows "Newest_on_path l a x y (find_newest_on_path l a x y)"
```

3.4.3 Explain

We can now define the explain operation, as described in the previous Subsection. In order to find the overall newest edge, we first compute the newest edge on the x branch,

then the one on the *y* branch, and then choose the larger one in a case distinction at the end.

TODO example

```
function (domintros) explain :: "ufe_data_structure ⇒ nat ⇒ nat ⇒ (nat *
    nat) set"
  where
    "explain ⦇uf_list = l, unions = u, au = a⦈ x y =
     (if x = y ∨ rep_of l x ≠ rep_of l y then {}
     else
        (let lca = lowest_common_ancestor l x y;
         newest_index_x = find_newest_on_path l a x lca;
         newest_index_y = find_newest_on_path l a y lca;
         (aₓ, bₓ) = u ! the (newest_index_x);
         (a_y, b_y) = u ! the (newest_index_y)
      in
       (if newest_index_x ≥ newest_index_y then
         {(aₓ, bₓ)} ∪ explain ⦇uf_list = l, unions = u, au = a⦈ x aₓ
           ∪ explain ⦇uf_list = l, unions = u, au = a⦈ bₓ y
       else
         {(a_y, b_y)} ∪ explain ⦇uf_list = l, unions = u, au = a⦈ x b_y
           ∪ explain ⦇uf_list = l, unions = u, au = a⦈ a_y y))
)"
  by pat_completeness auto
```

3.5 Proofs

This section introduces an invariant for the union find data structure and proves that the `explain` function terminates and is correct, when invoked with valid parameters.

3.5.1 Invariant and Induction Rule

The validity invariant of the data structure expresses that the data structure derived from subsequent unions with `ufe_union`, starting from the initial empty data structure. It also states that the unions were made with valid variables, i.e. variables which are in bounds.

```
abbreviation "ufe_invar ufe ≡
  valid_unions (unions ufe) (length (uf_list ufe)) ∧
  apply_unions (unions ufe) (initial_ufe (length (uf_list ufe))) = ufe"
```

With this definition, it is easy to show that the invariant holds after a union.

```
lemma union_ufe_invar:
  assumes "ufe_invar ufe"
  shows "ufe_invar (ufe_union ufe x y)"
```

It is also useful to prove that the old invariant, `ufa_invar`, is implied by the new invariant, so that we can use all the previously proved lemmas about `ufa_invar`. This is easily shown by computation induction on the function `apply_unions`, and by using the lemma from the Theory "Union Find", which states that `ufa_invar` holds after having applied `ufa_union`, and by proving that it holds for the initial empty data structure.

```
theorem ufe_invar_imp_ufa_invar: "ufe_invar ufe ⟹ ufa_invar (uf_list ufe)"
```

With this definition of the invariant, we can prove a new induction rule, which will be very useful for proving many properties of a union-find data structure. The induction rule, called `apply_unions_induct`, has as an assumption that the invariant holds for the given data structure *ufe*, and shows that a certain predicate holds for *ufe*. The base case that needs to be proven is that it holds for the initial data structure, and the induction step is that the property remains invariant after applying a union.

```
lemma apply_unions_induct[consumes 1, case_names initial union]:
  assumes "ufe_invar ufe"
  assumes "P (initial_ufe (length (uf_list ufe)))"
  assumes "⋀pufe x y. ufe_invar pufe ⟹ x < length (uf_list pufe) ⟹ y <
    length (uf_list pufe)
    ⟹ P pufe ⟹ P (ufe_union pufe x y)"
  shows "P ufe"
```

This induction rule can be used for most of the proofs about explain.

3.5.2 Termination Proof

An important result was to show that the function always terminates if `ufe_invar` holds. We will show this using `apply_unions_invar`, therefore we need to show that if the function terminates before `ufe_union` is applied, then it also terminates afterwards, assuming that *x* and *y* are in the same representative class, where *x* and *y* are the last two parameters of explain.

```
lemma explain_domain_ufe_union_invar:
  assumes "explain_dom (ufe, x, y)"
    and "ufe_invar ufe"
    and "rep_of (uf_list ufe) x = rep_of (uf_list ufe) y"
  shows "explain_dom (ufe_union ufe x2 y2, x, y)"
```

Proof. We can use the partial induction rule of explain, given that our first assumption is that explain terminates.

We show only the case when `newest_index_x` ≥ `newest_index_y`, because the other case is symmetric to it. The Isabelle code also contains proofs about the symmetry of explain, which are used in order to avoid duplicate proofs for the two cases of the

`explain` function, but they will not be discussed here, as they are not essential to prove the correctness of the function.

Initially, we remark that the lowest common ancestor and the newest index on path do not change after a union was applied. Therefore we will refer to the variables with the same names as in the function definition, e.g. *lca*, *ax*, etc., without specifying if we refer to e.g. `lowest_common_ancestor l x y` or `lowest_common_ancestor (ufa_union l x2 y2) x y`.

We assume that *x* and *y* are in the same representative class after the union. Given that (ax, bx) is the newest branch on the path fom *ax* to the lowest common ancestor *lca* of *x* and *y*, we know that every edge on the path from *x* to *ax* was also present before the union. Therefore $rep\_of(l, ax) = rep\_of(l, x)$ holds before the union, and we can apply the induction hypothesis and conclude `explain_dom(ufe_union ufe x2 y2, x, ax)`. (ax, bx) is also newer than the newest branch on the path fom *y* to the *lca*, therefore $rep\_of(l, y) = rep\_of(l, bx)$, and the induction hypothesis shows that `explain_dom(ufe_union ufe x2 y2, bx, y)`. The two recursive calls terminate, therefore `explain` terminates. $\qquad\square$

From this result we can prove the termination of `explain`:

```
theorem explain_domain:
  assumes "ufe_invar ufe"
  shows "explain_dom (ufe, x, y)"
```

Proof. We prove it by using `apply_union_induct`.

For the base case, we consider the empty data structure. There are no distinct variables with the same representative, therefore the algorithm terminates immediately.

For the induction step, if *x* and *y* are not in the same representative class after the union, the function terminates immediately. Otherwise, we can show that *x* and *ax* are in the same representative class before the union, and *bx* and *y* as well, therefore we can apply the previous lemma to the recursive calls of the function, and conclude that `explain` terminates. $\qquad\square$

3.5.3 Correctness Proof

There are two properties which define the correctness of explain: foremost, the equivalence closure of `explain x y` should contain the pair (x, y) (we shall refer to this property as "correctness"), additionally, the elements in the output should only be equations which are part of the input (we shall refer to this property as "validity"). The proposition about the validity of `explain` is the following:

```
theorem explain_valid:
  assumes "ufe_invar ufe"
    and "k ∈ (explain ufe x y)"
  shows "k ∈ set (unions ufe)"
```

We know from Subsection 3.5.2 that when the invariant holds, the function terminates. Therefore we can use the partial induction rule for `explain` that Isabelle automatically generates for partial functions. We can prove that *k* is a valid union, given that each element in `explain ufe x y` originally derives from the `unions` list, which is the list of input equations. In order to use this argument, we need to prove that the index found by `find_nearest_on_path` is in bounds.

```
lemma find_newest_on_path_Some:
  assumes "path l y p x"
    and "ufe_invar (|uf_list = l, unions = un, au = a|)"
    and "x ≠ y"
  obtains k where "find_newest_on_path l a x y = Some k ∧ k < length un"
```

This follows from the following lemma, that shows that the entries in the *associated unions* list are in bounds.

```
lemma au_valid:
  assumes "ufe_invar ufe"
    and "i < length (au ufe)"
  shows "au ufe ! i < Some (length (unions ufe))"
```

It is easily proven, given that all the values that are added to `au` by `ufe_union` are valid.

Thus we can prove the lemma about the validity of the explain function. It remains to show the correctness.

```
theorem explain_correct:
  assumes "ufe_invar ufe"
    and "rep_of (uf_list ufe) x = rep_of (uf_list ufe) y"
  shows "(x, y) ∈ (symcl (explain ufe x y))*"
```

Proof. This was shown using the induction rule of `explain`.

For the case where x = y, the algorithm returns the empty set, and because of reflexivity (x, y) is in the equivalence closure of the empty set.

As before, for the remaining cases we consider only the case where `newest_index_x` ≥ `newest_index_y`. From the induction hypothesis, we know that $(x, ax) \in$ (symcl (explain ufe x y))* and $(bx, y) \in$ (symcl (explain ufe x y))*.

Because of the definition of explain, it holds that $(ax, bx) \in$ (explain x y). Therefore from the transitivity of the equivalence closure it follows that $(x, y) \in$ (symcl (explain ufe x y))*. □

4 Congruence Closure with Explain Operation

4.1 Input equations

We now consider not only equations between constants, but also equations containing function symbols. Each function symbol is associated with an arity, which is the number of parameters it accepts. Function symbols with arity 0 are constants.

Arbitrary equations can be transformed to equations of depth at most 2, and with only one function of arity 2. This is done by currying and by introducing new constant symbols. See [9] for a detailed explanation of how this is done. In order to understand this thesis, it is irrelevant to know how the transformation is done, it is just important to know that the result is a set of equations between constant symbols and one specific function symbol F of arity 2. The congruence closure of these transformed equations is equal to the congruence closure of the original equations, therefore from this point on, we will only consider the transformed equations. Therefore there are two types of input equations for the congruence closure algorithm: either $a = b$ or $F(a, b) = c$ where a, b and c are constant symbols.

The data structure we use for these equations is the one described in Subsection 4.2.2.

4.2 Implementation

For the implementation of the congruence closure algorithm, we follow the description in the paper. [1] As before, the optimizations of path compression and considering the sizes of the representative classes are left out. These optimizations are not relevant for the correctness of the algorithm, and they could later be added to a refinement of the algorithm.

The algorithm uses a modified version of the union find algorithm for maintaining the equivalence classes between the constant elements, and it has a few additional data structures for storing equations containing the function symbol F. The following sections describe the modified union-find algorithm and then the congruence closure algorithm, as well as their correctness proofs.

4.2.1 Modified Union Find Algorithm

In order to implement an explain operation with a reasonable runtime for the congruence closure data structure, the paper [1] introduces an alternative union-find algorithm. Additionally to the union-find forest, there is a new data structure, the *proof forest*, i.e. a forest which has as nodes the variables, and as edges the unions that were made. Each time `ufa_union` is called on the union-find forest, the proof forest is modified with `add_edge`. In order to avoid the creation of cycles, redundant unions are ignored.

add_edge

The proof forest has directed edges, and for each equivalence class there is a representative node, where all the edges are directed towards. To keep this invariant, each time and edge from e to e' is added, all the edges on the path from the root of e to e are reversed. In the implementation, the proof forest is represented by a list which stores the parent of each node, exactly as in the union-find list. My implementation for adding an edge, which corresponds to the `union` operation, is the following:

```
function (domintros) add_edge :: "nat list ⇒ nat ⇒ nat ⇒ nat list"
  where
"add_edge pf e e' = (if pf ! e = e
                then (pf[e := e'])
                else add_edge (pf[e := e']) (pf ! e) e)"
  by pat_completeness auto
```

We can show that `add_edge e e'` terminates, if the invariant `ufa_invar` holds for the proof forest and e and e' do not belong to the same equivalence class.

```
lemma add_edge_domain:
  assumes "ufa_invar l" "rep_of l e ≠ rep_of l e'"
  shows "add_edge_dom (l, e, e')"
```

Proof. It can be proven by induction on the length of the path p from the root of e to e.

In the base case there is only one node in the path, therefore e must be equal to its root, therefore $pf!e = e$, and the algorithm terminates immediately.

In the other case e is not a root, then there is a path p' from the root to the parent of e which is shorter than the path from the root to e. The path p' is also present in the $pf[e := e']$, because the path does not contain e. Also, the representative of e in $pf[e := e']$ is equal to the representative of e', and the representative of the parent of e is still the old representative of e, therefore they are not in the same representative class, and we can apply the induction hypothesis and conclude that the recursive call terminates, therefore the function terminates. □

In order to prove the correctness of `add_edge`, we show that `ufa_union l x y` and `add_edge l x y` have the same behaviour, from which we can conclude that the equivalence classes of the union-find forest are the same as those of the proof forest. The theory `Union_Find` from the AFP[21] already provides the following lemma for `ufa_union`:

```
lemma ufa_union_aux:
  "rep_of (ufa_union l x y) i =
    (if rep_of l i = rep_of l x then rep_of l y else rep_of l i)"
```

We can show a similar lemma for `add_edge`:

```
lemma rep_of_add_edge_aux:
  assumes "rep_of l x ≠ rep_of l y"
  shows "rep_of (add_edge l x y) i =
    (if rep_of l i = rep_of l x then rep_of l y else rep_of l i)"
```

The additional assumption `rep_of l x ≠ rep_of l y` does not cause problems, because `add_edge` is only executed by the congruence closure algorithm when `rep_of l x ≠ rep_of l y`.

Proof. We already showed that the function terminates, therefore we can prove it by induction on `add_edge`.

The proof uses various lemmas about the behaviour of `rep_of` after a function update, which depends on where the function update was, and which element we want to find the representative of. These lemmas are proven by analysing how the paths in the forest change after a function update, which corresponds to adding a path in the forest. With these lemmas the induction is easily proven. □

We also show that `add_edge` has the expected behaviour, which is that it reverses all the edges from the root of e to e, and it adds an edge from e to e', i.e. after `add_edge` the forest contains a path from e' to the representative of e, which is the path which was there before `add_edge`, but reversed and with one added edge between e and e'. `rev` is a function which reverses a list, and `path_to_root` is the function described in Subsection 3.4.2. The proof can be shown by induction on `add_edge`.

```
lemma add_edge_correctness:
  assumes "ufa_invar pf"
    "rep_of pf e ≠ rep_of pf e'"
  shows "path (add_edge pf e e') e' ([e'] @ rev (path_to_root pf e)) (rep_of
    pf e)"
```

The proof forest has a similar structure as the union-find forest, therefore we prove that `add_edge` preserves the `ufa_invar` invariant from Section 3.2. This allows us to apply all the lemmas that were proven for the union-find forest also to the proof forest.

```
lemma add_edge_ufa_invar_invar:
  assumes "ufa_invar l"
    "rep_of l e ≠ rep_of l e'"
  shows "ufa_invar (add_edge l e e')"
```

Proof. In order to prove this lemma, we show another lemma which states that the invariant holds after a function update if the update does not cause the formation of a cycle. Then we show that each function update of add_edge does not form a cycle. □

add_label

Additionally, each edge is labeled with the input equation or the input equations which caused the adding of this edge. This is not necessary for the union-find algorithm by itself, but it will be needed by the explain operation for congruence closure. There are two possible types of labels: either an equation $a = b$ was input, or two equations of the type $F(a_1, a_2) = a$ and $F(b_1, b_2) = b$, where a_1 and b_1 were already in the same equivalence class before this union, as well as a_2 and b_2. In both theses cases a union between the equivalence classes of a and b must be made. The labeling is implemented by using an additional list, which at each index contains the label of the outgoing edge, or None if there is no outgoing edge. It is similar to the associated unions list of union-find, but it contains directly the labels instead of an index to another list.

The labels have the type pending_equation, which can be either one or two equations.

```
datatype pending_equation = One equation
  | Two equation equation
```

The name pending_equation derives from the fact that it is also the type of the elements of the pending list, which will be described in the next section. Theoretically this allows also for invalid equations for example two equations of the type $a = b$ and $c = d$, but we will prove in the next sections that the equations in the labels list are always either One ($a = b$) or Two ($F(a_1, a_2) = a$) ($F(b_1, b_2) = b$).

Each time an edge gets added to the proof forest, the labels need to be updated as well. The function add_label adds a label to the new edge, and modifies the labels for the edges which are modified by add_edge:

```
function (domintros) add_label :: "pending_equation option list ⇒ nat list
    ⇒ nat ⇒ pending_equation ⇒ pending_equation option list"
  where
"add_label pfl pf e lbl =
   (if pf ! e = e
      then (pfl[e := Some lbl])
      else add_label (pfl[e := Some lbl]) pf (pf ! e) (the (pfl ! e)))"
  by pat_completeness auto
```

Similarly to the `path_to_root` function, `add_label` has the same recursive calls as `rep_of`, therefore it has the same domain.

```
lemma rep_of_dom_iff_add_label_dom:
  "rep_of_dom (pf, y) ⟷ add_label_dom (pfl, pf, y, y')"
```

4.2.2 Congruence Closure Data Structure

For the congruence closure algorithm there are five important data structures, which are described in the following. More details on this topic can be found in [1].

- `cc_list`: the union-find list, corresponds to the `uf_list`.

- `use_list`: a two-dimensional list which contains for each representative a a list of input equations $F(b_1, b_2) = b$ where the representative of b_1 or b_2 is a.

- `lookup`: a lookup table indexed by pairs of representatives b and c, which stores an input equation $F(a_1, a_2) = a$ such that b is the representative of a_1 and c is the representative of a_2, or `None` if no such equation exists.

- `pending`: equations of the type `One` $(a = b)$ or `Two` $(F(a_1, a_2) = a)$ $(F(b_1, b_2) = b)$ where a and b need to be merged, and a_1 and b_1 are already in the same congruence class, as well as a_2 and b_2.

- `proof_forest`: the proof forest as described in the previous subsection.

- `pf_labels`: the labels of the proof forest as described in the previous subsection.

- `input`: a set of the input equation, which will be useful for some proofs in the next sections.

In the following, we shall refer to `cc_list` as l, the use list as u, the lookup table as t, the pending list as pe, the proof forest as pf, the labels list for the proof forest as pfl and the input as ip, unless otherwise stated.

4.2.3 Congruence Closure Algorithm

With this data structure we can implement the `merge` function as described in [1]. It takes as parameter the current congruence closure data structure and an equation which it adds to the data structure. It uses the `propagate` function, which will be described later, which performs unions between the constant symbols in the pending list.

If the input equation is of the type $F(a_1, a_2) = a$, then there are two possibilities: if there is already an equation $F(b_1, b_2) = b$ in the lookup table at the index

$(rep\_of(l, a_1), rep\_of(l, a_2))$, then we know that $a_1 = b_1$ and $a_2 = b_2$, and we add $F(b_1, b_2) = b$ and $F(a_1, a_2) = a$ to pending, so that the equivalence classes of a and b will be merged. On the other hand, if the respective lookup entry is None, then the equation is added to the lookup table, at the index $(rep\_of(l, a_1), rep\_of(l, a_2))$ so that the next time an equation with congruent parameters is input, they will be added together to pending.

For this case distinction there is a function `lookup_Some`, which returns `True` if there is an entry in lookup at the index $(rep\_of(l, a_1), rep\_of(l, a_2))$ and False otherwise, and a function `update_lookup`, which adds the equation to lookup at the index $(rep\_of(l, a_1), rep\_of(l, a_2))$.

```
fun merge :: "congruence_closure ⇒ equation ⇒ congruence_closure"
 where
"merge ⦇cc_list = l, use_list = u, lookup = t, pending = pe, proof_forest =
    pf, pf_labels = pfl, input = ip⦈
(a ≈ b) =
 propagate
   ⦇cc_list = l, use_list = u, lookup = t, pending = One (a ≈ b)#pe,
       proof_forest = pf, pf_labels = pfl, input = insert (a ≈ b) ip⦈"

| "merge ⦇cc_list = l, use_list = u, lookup = t, pending = pe, proof_forest =
     pf, pf_labels = pfl, input = ip⦈
(F a₁ a₂ ≈ a) =
(if (lookup_Some t l (F a₁ a₂ ≈ a))
 then propagate ⦇cc_list = l, use_list = u, lookup = t,
         pending = link_to_lookup t l (F a₁ a₂ ≈ a)#pe, proof_forest = pf,
            pf_labels = pfl, input = insert (F a₁ a₂ ≈ a) ip⦈
 else ⦇cc_list = l,
       use_list = (u[rep_of l a₁ := (F a₁ a₂ ≈ a)#(u ! rep_of l a₁)])[rep_of
          l a₂ := (F a₁ a₂ ≈ a)#(u ! rep_of l a₂)],
       lookup = update_lookup t l (F a₁ a₂ ≈ a),
       pending = pe, proof_forest = pf, pf_labels = pfl, input = insert (F a₁
          a₂ ≈ a) ip⦈
)"
```

The main part of the algorithm is executed in `propagate`, which recursively takes one item from pending and performs the union of the representative classes. As previously mentioned, the pending item could be either an equation of the type $a = b$, or two equations of the type $F(a_1, a_2) = a$ and $F(b_1, b_2) = b$, where a_1 and a_2 are already in the same representative class as b_1 and b_2 respectively. In both cases the representative classes of a and b need to be merged. The functions `left` and `right` simply retrieve a and b from either of the two types of pending equations. If a and b are already in the same representative class, nothing needs to be done, otherwise the union is performed. For more clarity, the union is defined separately as `propagate_step`.

```
function propagate :: "congruence_closure ⇒ congruence_closure"
  where
"propagate ⦇cc_list = l, use_list = u, lookup = t, pending = [], proof_forest
    = pf, pf_labels = pfl, input = ip⦈ =
⦇cc_list = l, use_list = u, lookup = t, pending = [], proof_forest = pf,
    pf_labels = pfl, input = ip⦈"
| "propagate
⦇cc_list = l, use_list = u, lookup = t, pending = (eq # pe), proof_forest =
    pf, pf_labels = pfl, input = ip⦈ =
(let a = left eq; b = right eq in
  (if rep_of l a = rep_of l b
    then propagate ⦇cc_list = l, use_list = u, lookup = t, pending = pe,
        proof_forest = pf, pf_labels = pfl, input = ip⦈
    else
      propagate (propagate_step l u t pe pf pfl ip a b eq)
))"
  by pat_completeness auto
```

The union consists of the previously discussed `ufa_union`, `add_edge` and `add_label`, as well as a loop which moves all elements from the use list of $rep\_of(l,a)$ to either $rep\_of(l,b)$, or to pending. This is necessary, because the old representative of a is not a representative any more, and its new representative is $rep\_of(l,b)$.

```
abbreviation propagate_step
  where
"propagate_step l u t pe pf pfl ip a b eq ≡
  propagate_loop (rep_of l b) (u ! rep_of l a)
    ⦇cc_list = ufa_union l a b,
    use_list = u[rep_of l a := []],
    lookup = t,
    pending = pe,
    proof_forest = add_edge pf a b,
    pf_labels = add_label pfl pf a eq,
    input = ip⦈"
```

The loop is defined as a recursive function, which considers each element of the use list of $rep\_of(l,a)$, and either adds it to the use list and the lookup table, or if there is already an entry in lookup, then that entry together with the current equation are added to pending.

```
fun propagate_loop
  where
"propagate_loop rep_b (u1 # urest)
⦇cc_list = l, use_list = u, lookup = t, pending = pe, proof_forest = pf,
    pf_labels = pfl, input = ip⦈
=
  propagate_loop rep_b urest (
    if (lookup_Some t l u1)
```

```
    then
      ⦇cc_list = l, use_list = u, lookup = t,
          pending = link_to_lookup t l u1#pe,
          proof_forest = pf, pf_labels = pfl, input = ip⦈
    else
      ⦇cc_list = l,
          use_list = u[rep_b := u1 # (u ! rep_b)],
          lookup = update_lookup t l u1,
          pending = pe, proof_forest = pf, pf_labels = pfl, input = ip⦈
)"
| "propagate_loop _ [] cc = cc"
```

The function `are_congruent` returns `True` if an equation is in the congruence closure of all the input equations so far. It simply checks if the elements have the same representative or if they have the same representative as the corresponding entry in lookup.

```
fun are_congruent :: "congruence_closure ⇒ equation ⇒ bool"
  where
"are_congruent ⦇cc_list = l, use_list = u, lookup = t, pending = pe,
    proof_forest = pf, pf_labels = pfl, input = ip⦈ (a ≈ b) =
   (rep_of l a = rep_of l b)"
| "are_congruent ⦇cc_list = l, use_list = u, lookup = t, pending = pe,
    proof_forest = pf, pf_labels = pfl, input = ip⦈ (F a₁ a₂ ≈ a) =
   (case lookup_entry t l a₁ a₂ of
     Some (F b₁ b₂ ≈ b) ⇒ (rep_of l a = rep_of l b)
   | None ⇒ False
)"
```

4.3 Correctness Proof

4.3.1 Invariants

At this point we can already prove some properties of the congruence closure data structure. Our approach this time is different than the one for the union-find algorithm. Instead of defining an induction rule like in the union-find section and then prove the properties through the induction rule, we define the properties as invariants and then prove that they remain invariant after applying merge. For each invariant, we need to follow the same steps. They are listed here in order to introduce the name for each step:

1. Prove that the invariant holds for the initial empty congruence closure.

2. Prove that if the invariant holds before the merge operation, it also holds after the merge operation. Below is a list of what needs to be proven:

a) The invariant holds after one step in the `propagate_loop`. We shall refer to the two cases of the function as `loop1` and `loop2`.

b) The invariant holds after the entire `propagate_loop`.

c) The invariant holds for the parameters of `propagate_loop` in `propagate_step`. We shall refer to the this case as `mini_step`.

d) It holds after `propagate_step`.

e) It holds after `propagate`.

f) And finally, it holds after `merge`.

Let us now look at the concrete invariants. Each list in the data structure has an invariant which states that all the elements which are in the list are in bounds. The corresponding proofs are easy to prove if we assume that all the input equations contain only valid elements.

One of the invariants is the usual `ufa_invar` that we know from the union-find algorithm. The `ufa_invar` holds for the `cc_list` and the `proof_forest`. These two are only modified before entering the in the `mini_step`, and we already proved previously that the `ufa_invar` holds after `ufa_union` (in Section 3.2) and after `add_edge` (Subsection 4.2.1). Therefore it also holds after `merge`.

We define a new invariant `inv_same_rep_classes`, which states, as the name suggests, that the union-find forest and the proof forest represent the same equivalence classes:

`rep_of l i = rep_of l j ⟷ rep_of pf i = rep_of pf j`

This invariant is important for the proofs that consider `add_edge`. That is because we only showed that `add_edge pf x y` terminates if `rep_of pf x ≠ rep_of pf y`. `propagate` only executes `add_edge` when `rep_of l x ≠ rep_of l y`. The invariant shows that these two statements are equivalent, therefore `add_edge` always terminates when used inside of the algorithm.

In order to prove the invariant, given that the two lists l and pf are only modified during the `mini_step`, it is sufficient to show that `ufa_union l x y` and `add_edge pf x y` have the same behaviour, which is what we showed in Subsection 4.2.1.

Furthermore, for each data structure there is an invariant which states the properties which were informally described in Subsection 4.2.2. They are described in the following.

For the use list, the invariant `use_list_invar` states that for each representative a, its use list only contains equations of the type $F(b_1, b_2) = b$, where a is the representative of either b_1 or b_2.

Proof. For the correctness proof after the `propagate_loop`, we need to add an additional assumption that the second parameter only contains equations of the type $F(a_1, a_2) = a$

and the representative of either a_1 or a_2 is $rep\_of(l,b)$ (where b is the right side of the equation which is being propagated). This follows from the facts that the second parameter of `propagate_loop` is $(u!rep\_of(l,a))$, the invariant holds before the `propagate_loop`, and the new representative of a after the union is b.

With this assumption we can show that each time an equation gets added to the use list in the `propagate_loop`, it is a valid equation.

In the proof after the `merge` operation, the use list is only modified in the third case, and only equations of the form $F(a_1, a_2) = a$ are added to $rep\_of(l, a_1)$ and $rep\_of(l, a_2)$. Therefore all the necessary properties hold for these new equations.

For the remaining cases, use list is either unchanged, or something is removed from it, therefore the invariant trivially holds. □

The invariant `lookup_invar` for lookup is similar, it states that each entry in the lookup table at index (i, j), for representatives i and j, is either `None` or is an equation of the form $F(a_1, a_2) = a$ where the representative of a_1 is i and the representative of a_2 is j.

Proof. Each time an equation is added to lookup, it has the desired form and it is added to the index $(rep\_of(l, a_1), rep\_of(l, a_2))$. This happens in the `propagate_loop` and in `merge`. In the `propagate_loop`, the added equation derives from the use list, for which we proved with the previous invariant that its equations have the desired form. In `merge`, only the equations of the type $F(a_1, a_2) = a$ are added to lookup. □

For pending, the invariant `pending_invar` states that the equations are either of the form `One` $(a = b)$ or `Two` $(F(a_1, a_2) = a)$ $(F(b_1, b_2) = b)$ where $rep\_of(l, a_1) = rep\_of(l, b_1)$ and $rep\_of(l, a_2) = rep\_of(l, b_2)$:

Proof. We need to show that in the `propagate_loop` the equation $u1$ we add to pending has a valid form. We know that $u1$ derives from the use list, therefore it is of the form $F(a_1, a_2) = a$. Then we link to it the lookup entry at the index $(rep\_of(l, a_1), rep\_of(l, a_2))$. From the lookup invariant we know that there is an entry of the form $F(b_1, b_2) = b$ at this index where $rep\_of(l, a_1) = rep\_of(l, b_1)$ and $rep\_of(l, a_2) = rep\_of(l, b_2)$. This shows that they are valid equations for pending.

The same holds for the equations added to pending in `merge`. □

There is also an invariant which states that the `cc_list`, the first dimension of the use list, both dimensions of lookup, the proof forest and the `pf_labels` have the same length. This was trivial to prove, given that the algorithm never changes the length of the lists, and initially the lists have the same length.

The remaining invariants will be described later on, when they become relevant.

4.3.2 Abstract Formalization of Congruence Closure

In order to prove the correctness of the algorithm, we define an abstraction of congruence closure. We cannot use any previously defined definitions, because the data structure that we use can only represent a subset of all possible equations, for example it cannot represent equations of the type $a = F(b, c)$ or $F(F(a, b), c) = d$. For this reason, we define an inductive set which represents the congruence closure of a set of equations and only uses our restricted definition of equation.

```
inductive_set Congruence_Closure :: "equation set ⇒ equation set" for S
  where
    base: "eqt ∈ S ⟹ eqt ∈ Congruence_Closure S"
  | reflexive: "(a ≈ a) ∈ Congruence_Closure S"
  | symmetric: "(a ≈ b) ∈ Congruence_Closure S ⟹ (b ≈ a) ∈
      Congruence_Closure S"
  | transitive1: "(a ≈ b) ∈ Congruence_Closure S ⟹ (b ≈ c) ∈
      Congruence_Closure S
⟹ (a ≈ c) ∈ Congruence_Closure S"
  | transitive2: "(F a₁ a₂ ≈ b) ∈ Congruence_Closure S ⟹ (b ≈ c) ∈
      Congruence_Closure S
⟹ (F a₁ a₂ ≈ c) ∈ Congruence_Closure S"
  | transitive3: "(F a₁ a₂ ≈ a) ∈ Congruence_Closure S
⟹ (a₁ ≈ b₁) ∈ Congruence_Closure S ⟹ (a₂ ≈ b₂) ∈ Congruence_Closure S
⟹ (F b₁ b₂ ≈ a) ∈ Congruence_Closure S"
  | monotonic: "(F a₁ a₂ ≈ a) ∈ Congruence_Closure S ⟹ (F a₁ a₂ ≈ b) ∈
      Congruence_Closure S
⟹ (a ≈ b) ∈ Congruence_Closure S"
```

The following proof rule follows directly from the definition of congruence closure, and proved to be very useful for multiple proofs:

```
lemma Congruence_Closure_eq[case_names left right]:
  assumes "⋀ a. a ∈ A ⟹ a ∈ Congruence_Closure B"
    "⋀ b. b ∈ B ⟹ b ∈ Congruence_Closure A"
  shows "Congruence_Closure A = Congruence_Closure B"
```

It is used to prove equality between congruence closures of *A* and *B*. It states that it is sufficient to prove that all elements of set *A* are in the congruence closure of *B* and vice versa, instead of having to prove that all elements of the congruence closure of *A* are in the congruence closure of *B*.

4.3.3 Correctness

To prove the correctness of the congruence closure implementation, we need to show that the invariants imply that `are_congruent cc eq` returns `True` if and only if the equation *eq* lies in the congruence closure of the input equations.

```
theorem are_congruenct_correct:
  assumes "cc_invar cc" "pending cc = []"
  shows "eq ∈ Congruence_Closure ((input cc)) ⟷ are_congruent cc eq"
```

The paper [1] proves this by stating the folllowing invariant which holds throughout the algorithm,

```
Congruence_Closure(representativeE ∪ pending) = Congruence_Closure (input)
```

where `representativeE` can be seen as the set of equations derived from our union-find list and the equations in lookup. It is the union of the following two sets:

- `representative_set` is defined such that its congruence closure contains all the equations between two elements which have the same representative.

- `lookup_entries_set` is the set of all the entries in lookup at indexes which are representatives.

```
abbreviation representatives_set :: "nat list ⇒ equation set"
  where
    "representatives_set l ≡ {a ≈ rep_of l a |a. l ! a ≠ a}"
```

```
abbreviation lookup_entries_set :: "congruence_closure ⇒ equation set"
  where
    "lookup_entries_set cc ≡ {F a' b' ≈ rep_of (cc_list cc) c | a' b' c c₁ c₂ .

                cc_list cc ! a' = a' ∧ cc_list cc ! b' = b'
                ∧ lookup cc ! a' ! b' = Some (F c₁ c₂ ≈ c)}"
```

```
definition representativeE :: "congruence_closure ⇒ equation set"
  where
    "representativeE cc = representatives_set (cc_list cc) ∪
        lookup_entries_set cc"
```

The formal definition of the aforementioned invariant is the following, where `pending_set` converts the pending list to a set of equations of the type $a = b$:

```
definition inv2 :: "congruence_closure ⇒ bool"
  where
    "inv2 cc ≡
Congruence_Closure (representativeE cc ∪ pending_set (pending cc)) =
    Congruence_Closure (input cc)"
```

The set of input equations is only modified by the `merge` function, but remains constant throughout the `propagate` function, therefore for the proof we just need to show that the congruence closure of the `representativeE` set and pending remain unchanged after the `propagate` function.

The main challenge is to prove that the invariant holds after the `mini_step`. We will show that `Congruence Closure (representativeE ∪ pending)` before the `propagate_step` is equal to `Congruence Closure (representativeE ∪ pending ∪ (u ! rep_of l a))` after the `mini_step`. Then we prove that the latter is equal to `Congruence Closure (representativeE ∪ pending)` after the `propagate_loop`. These two lemmas imply that the congruence closure of the `representativeE` set and pending remain unchanged after the `propagate` function.

We will first prove the second statement, given that it is much easer to show.

Proof. We need to show that `Congruence Closure (representativeE ∪ pending ∪ (u ! rep_of l a))` before the `propagate_loop` is equal to `Congruence Closure (representativeE ∪ pending)` after the `propagate_loop`. In each step of the loop, one element from $(u!rep\_of(l,a))$ is moved either to pending or to lookup. Therefore after the loop each element of $(u!rep\_of(l,a))$ is either in pending or in representativeE. The elements of pending and representativeE are never removed from the set, therefore they are present also after the loop. □

The first lemma is more difficult to prove. The following is the statement of the lemma:

```
lemma inv2_mini_step:
  assumes "a = left eq" "b = right eq"
  "cc_invar ⦇cc_list = l, use_list = u, lookup = t, pending = (eq # pe),
  proof_forest = pf, pf_labels = pfl, input = ip⦈"
    shows "Congruence_Closure
(representativeE
⦇cc_list = l, use_list = u, lookup = t, pending = (eq # pe),
proof_forest = pf, pf_labels = pfl, input = ip⦈
∪ pending_set (eq # pe))
=
Congruence_Closure (representativeE
⦇cc_list = ufa_union l a b,
   use_list = u[rep_of l a := []],
   lookup = t,
   pending = pe,
   proof_forest = add_edge pf a b,
   pf_labels = add_label pfl pf a eq,
   input = ip⦈
∪ pending_set pe
∪ set (u ! rep_of l a))"
```

Proof. There are two inclusions which need to be shown. We can use the rule `Congruence_Closure_eq` from Subsection 4.3.2, which means that it is sufficient to show that each equation in

the set on the left-hand side is in the congruence closure of the right-hand side and vice versa.

"⊆" It needs to be shown that the equations of the `representatives_set`, in lookup and in pending are in the Congruence Closure of the right-hand side.

Regarding the `representatives_set`, all the elements which had the same representative before a union also have the same representative after a union.

For the pending set, we need to prove that the equation that is removed from pending is still in the congruence closure after the `mini_step`. This holds, because the equation which is removed is $a = b$, and a and b are in the same equivalence class after the `ufe_union`.

The problematic cases are the equations in lookup. Given that after the union $rep\_of(l,a)$ is not a root any more, the entries in lookup which have as first or second index $rep\_of(l,a)$ are not in the `lookup_set` anymore after the union. The goal is to prove that these equations are exactly the equations which are present in $(u!rep\_of(l,a))$, but until now it was only proven that the equations in the use list are valid, not that they are exhaustive. A new invariant `use_list_inv2` is needed which states that all elements which are present in the lookup table at index (i,j) are also in the corresponding use lists of i and j. We will introduce this invariant later.

"⊇" We need to show that the equations of the `representatives_set`, `lookup_entries_set`, pending set and of $(u!rep\_of(l,a))$ on the right-hand side are in the congruence closure of the left-hand side.

The `representatives_set` contains equations of the type $c = rep\_of(ufa\_union(l,a,b),c)$ for each element c which is not a root. If the representative after the union is the same before the union, the same equation is in the `representatives_set` of the left-hand side. The only representative that is different than before the union is the representative of a, which has as new representative $rep\_of(l,b)$. The left-hand side contains the equations $b = rep\_of(l,b)$ and $a = b$ (which is in pending). By transitivity, the congruence closure also contains $a = rep\_of(l,b)$ and $rep\_of(l,b)$ is exactly the same as $rep\_of(ufa\_union(l,a,b),a)$.

Regarding lookup, all the elements which are roots after the union, are also roots before the union, therefore all elements in the `lookup_entry_set` of the right-hand side are also in the left-hand side.

It is evident that the equations in pending on the right-hand side are also in pending in the left-hand side.

It is more difficult to show that the equations in $(u!rep\_of(l,a))$ are also present in the lookup table of the left-hand side. As before, we need a new invariant `lookup_invar2`, which states that all equations which are in the use list of a root i are also present in the lookup table. Below is the description of the new invariants. □

We need two new invariants of this form:

- `lookup_invar2`: The elements in the lookup table are also present in the use list.

- `use_list_invar2`: The elements in the use list are also present in the lookup table.

Unfortunately, these two invariants are not exactly true, because if there are two different equations where the elements have the same representatives, then they cannot both be present in the lookup table, because it only stores one equation for each pair of representatives. In fact, the set of equations in lookup and in the use list are not exactly the same, but for each equation in one of them, there is a "similar" equation in the other one.

The difficulty was to find a suitable definition of "similar" which is not too strong, otherwise it wouldn't be true, but also not too weak, otherwise it is not possible to prove the invariant `inv2`.

The right definition of similar turned out to be the following:

Definition. Two equations $F(a_1, a_2) = a$ and $F(b_1, b_2) = b$ are *similar*, if $rep\_of(l, a_1) = rep\_of(l, b_1)$, $rep\_of(l, a_2) = rep\_of(l, b_2)$ and $(a = b) \in Congruence\_Closure(representatives\_set \cup pending)$.

Simply stating that a and b have the same representative would be too strong, because during the propagate function, they are added to pending in order to be merged later, and are not merged yet. If we use `representativeE` instead of `representatives_set`, the invariant is not strong enough in order to prove `inv2_mini_step`.

The final invariants are the following:

- `lookup_invar2`: For each equation in lookup at the index (i, j) (where i and j are representatives) there is a similar equation in use list i and one in use list j.

- `use_list_invar2`: For each equation $F(c_1, c_2) = c$ in use list at the index i (where i is a representative) there is a similar equation in lookup at the index $(rep\_of(l, c_1), rep\_of(l, c_2))$.

Here follows the proof for `lookup_invar2`:

Proof. We need to show that if the invariant holds before `merge`, then it also holds after the `merge`.

The main aim is to show that it holds after `propagate`. We assume that before `propagate` the invariants `lookup_invar2` and `use_list_invar2` hold. In particular, we observe that for each equation $F(c_1, c_2) = c$ in $(u!rep\_of(l, a))$ (which we shall refer to with u_a) there is a similar equation in lookup at the index $(rep\_of(l, c_1), rep\_of(l, c_2))$ and therefore there are similar equations in $(u!rep\_of(l, c_1))$ and in $(u!rep\_of(l, c_2))$.

In the `propagate_loop` u_a is emptied, while the other use lists are not modified, and u_a is handed over as a parameter to the `propagate_loop`.

From now on let l be the `cc_list` after the `ufe_union`.

In `loop1` the lookup table and the use lists are not modified, thus there is nothing to show.

In `loop2` we take an equation $u1$ of the form $F(c_1, c_2) = c$ from u_a.

$u1$ is then added to lookup at the index $(rep\_of(l, c_1), rep\_of(l, c_2))$. We need to show that after this step, an equation similar to $u1$ is present both in the use list of $rep\_of(l, c_1)$ and the use list of $rep\_of(l, c_2)$.

This holds if none of those two use lists are u_a, because of the observation made earlier, and because no use list has not been modified apart from u_a.

If one (or both) of the use lists are u_a, then the representative of the corresponding element was $rep\_of(l, a)$ before the union, therefore after the union it is $rep\_of(l, b)$. Given that $u1$ is also added to $u!rep\_of(l, b)$ by the function, we can conclude that there is a similar equation also in this use list. $\qquad\square$

Here follows the proof for `use_list_invar2`:

Proof. We need to show that if before the `merge` the invariant holds, then it also holds after the `merge`.

The difficulty in this proof was when after the `mini_step`, because of the union, $rep\_of(l, a)$ is not a root any more, therefore if there are some equations $F(c_1, c_2) = c$ in the use list where the representative of c_1 or c_2 is $rep\_of(l, a)$, they have a new representative after the union, which means that the corresponding similar equations in the lookup table are not at the right index anymore.

If $F(c_1, c_2) = c$ was in $u!rep\_of(l, a)$ (which we shall again refer to with u_a), then it is removed from the use list after the `mini_step`, and the previous remark does not cause any problems. However, there could be equations in use list at an index which is not $rep\_of(l, a)$, where c_1 or c_2 has the same representative of a. These equations do not have a similar equation in lookup after the `mini_step`, but we know from `lookup_invar2` that they have a similar equation in u_a. Therefore, in order show that `use_list_invar2` holds after the `propagate_loop`, it is sufficient to show that for each equation in u_a a similar equation is added (or already present) in the lookup table after the `propagate_loop`.

From now on let l be the `cc_list` after the `ufe_union`.

The `propagate_loop` removes an equation $F(c_1, c_2) = c$ from u_a, and it enters `loop1` when lookup contains an equation $F(d_1, d_2) = d$ at the index $(rep\_of(l, c_1), rep\_of(l, c_2))$. Then the equation $c = d$ is added to pending. Note that $F(c_1, c_2) = c$ and $F(d_1, d_2) = d$

are similar at this point, because $rep\_of(l, c_1) = rep\_of(l, d_1)$ and $rep\_of(l, c_2) = rep\_of(l, d_2)$ follows from the lookup_invar, and $c = d$ is added to pending.

This case is exactly the reason why we can only prove that there is a "similar" equation in lookup, and not exactly the same.

Regarding loop2, it is entered when lookup contains None at the index $(rep\_of(l, c_1), rep\_of(l, c_2))$. Then $F(c_1, c_2) = c$ is added to $(u!rep\_of(l, b))$ and to lookup. Obviously, an equation is similar to itself, therfore after this lookup contains a similar equation to $F(c_1, c_2) = c$.

There is also a new element which is added to the use list, and it has a similar equation in lookup, which is $F(c_1, c_2) = c$ itself, therefore the invariant holds. □

With these two invariants the proof for inv2 is completed. Given that the pending list is always empty after the termination of propagate, the correctness proof are_congruent_correct, which was defined at the beginning of this section, follows directly from this invariant.

All the invariants are put together in the invariant cc_invar, and using all the previously described proofs about the invariant, we can prove that cc_invar holds for the initial empty data structure, and that it holds after merge.

theorem cc_invar_initial_cc: "cc_invar (initial_cc n)"

Proof. All the above-mentioned invariants hold trivially for the initial case, given that all the data structures are empty or contain only None in the beginning. □

theorem cc_invar_merge:
  assumes "cc_invar cc" "valid_vars eq (nr_vars cc)"
  shows "cc_invar (merge cc eq)"

Proof. We already proved for each individual invariant, that they hold after propagate. The proof for merge uses the fact that propagate terminates, which will be proven in the following section. □

4.3.4 Termination

We already proved that the functions add_edge and add_label terminate, the only missing proof is the termination of propagate. All the remaining functions are simple enough for Isabelle to prove their termination automatically.

In order to prove the termination of propagate, we show that the number of equivalence classes strictly decreases in each step of propagate, therefore the function terminates at the latest when all the elements belong to the same equivalence class.

The number of equivalence classes is defined as the number of roots in the union-find forest. The function card returns the cardinality of a set.

```
abbreviation root_set
  where
    "root_set l ≡ {i | i. i < length l ∧ l ! i = i}"

definition nr_eq_classes :: "nat list ⇒ nat"
  where
    "nr_eq_classes l = card (root_set l)"
```

With this, we can show that after a union, $rep\_of(l, a)$ is not a root anymore, therefore there is one less root in the forest.

```
lemma ufa_union_decreases_nr_eq_classes:
  assumes "ufa_invar l" "a < length l"
    "rep_of l a ≠ rep_of l b"
  shows "nr_eq_classes (ufa_union l a b) = nr_eq_classes l - 1"
```

The termination proof for `propagate` follows from this lemma, with an additional assumption that there is at least one variable, so that there is at least one equivalence class.

```
lemma propagate_domain:
  assumes "cc_invar cc" "nr_vars cc > 0"
  shows "propagate_dom cc"
```

Proof. We prove it by induction on the amount of equivalence classes in the union-find forest.

If the pending list is empty, the function terminates. Otherwise, the first element *eq* is taken from the pending list. We define *a* as `left eq` and *b* as `right eq`, as in the function.

If *a* and *b* are already in the same equivalence class, then the union-find list is not modified, therefore we cannot use the induction hypothesis. Therefore we prove it by induction on the length of the pending list.

If *a* and *b* are not in the same equivalence class, they are merged by the `propagate_step`, therefore the number of equivalence classes descreases by one according to the lemma `ufa_union_decreases_nr_eq_classes` and we can prove the goal by using the induction hypothesis. □

4.4 Implementation of the Explain Operation

I implemented the explain operation for congruence closure, leaving the proof of termination and correctness open for future work. Nevertheless I would like to dedicate this section on the description of the implementation and a proposal of how the correctness could be proven.

TODO

5 Conclusion

Correctness proof complicated because of many invariants about validity, in bounds, right structure, etc.

Termination proof easier, no invariants

5.1 Future work

List of Figures

Bibliography

[1] R. Nieuwenhuis and A. Oliveras. "Proof-Producing Congruence Closure." In: *Lecture Notes in Computer Science*. Springer Berlin Heidelberg, 2005, pp. 453–468. DOI: `10.1007/978-3-540-32033-3_33`.

[2] P. Lammich. "Refinement to Imperative HOL." In: *Journal of Automated Reasoning* 62.4 (Oct. 2017), pp. 481–503. DOI: `10.1007/s10817-017-9437-1`.

[3] B. A. Galler and M. J. Fisher. "An improved equivalence algorithm." In: *Communications of the ACM* 7.5 (May 1964), pp. 301–303. DOI: `10.1145/364099.364331`.

[4] R. E. Tarjan. "A class of algorithms which require nonlinear time to maintain disjoint sets." In: *Journal of Computer and System Sciences* 18.2 (Apr. 1979), pp. 110–127. DOI: `10.1016/0022-0000(79)90042-4`.

[5] R. E. Shostak. "An algorithm for reasoning about equality." In: *Communications of the ACM* 21.7 (July 1978), pp. 583–585. DOI: `10.1145/359545.359570`.

[6] G. Nelson and D. C. Oppen. "Fast Decision Procedures Based on Congruence Closure." In: *Journal of the ACM* 27.2 (Apr. 1980), pp. 356–364. DOI: `10.1145/322186.322198`.

[7] P. J. Downey, R. Sethi, and R. E. Tarjan. "Variations on the Common Subexpression Problem." In: *Journal of the ACM* 27.4 (Oct. 1980), pp. 758–771. DOI: `10.1145/322217.322228`.

[8] L. de Moura and N. Bjørner. "Z3: An Efficient SMT Solver." In: *Tools and Algorithms for the Construction and Analysis of Systems*. Springer Berlin Heidelberg, 2008, pp. 337–340. DOI: `10.1007/978-3-540-78800-3_24`.

[9] R. Nieuwenhuis and A. Oliveras. "Fast congruence closure and extensions." In: *Information and Computation* 205.4 (Apr. 2007), pp. 557–580. DOI: `10.1016/j.ic.2006.08.009`.

[10] T. Nipkow, L. C. Paulson, and M. Wenzel. *Isabelle/HOL A Proof Assistant for Higher-Order Logic. A Proof Assistant for Higher-Order Logic*. Springer London, Limited, 2003, p. 226. ISBN: 9783540459491.

[11] A. Krauss. "Defining Recursive Functions in Isabelle/HOL." In: (May 2012).

[12] P. Lammich. "Collections Framework." In: *Archive of Formal Proofs* (Dec. 2009). `https://isa-afp.org/entries/Collections.html,` Formal proof development. ISSN: 2150-914x.

[13] S. Conchon and J.-C. Filliâtre. "A persistent union-find data structure." In: *Proceedings of the 2007 workshop on Workshop on ML - ML '07*. ACM Press, 2007. DOI: `10.1145/1292535.1292541`.

[14] M. P. L. Haslbeck and P. Lammich. "Refinement with Time - Refining the Run-Time of Algorithms in Isabelle/HOL." en. In: (2019). DOI: `10.4230/LIPICS.ITP.2019.20`.

[15] P. Corbineau. "Autour de la clôture de congruence avec Coq." MA thesis. Université Paris-Sud, 2001.

[16] *congruence tactic in the Coq documentation*. `https://coq.inria.fr/distrib/V8.11.2/refman/proof-engine/tactics.html#coq,` Accessed 25.07.2022.

[17] D. Selsam and L. de Moura. "Congruence Closure in Intensional Type Theory." In: *Automated Reasoning*. Springer International Publishing, 2016, pp. 99–115. DOI: `10.1007/978-3-319-40229-1_8`.

[18] J. Blanchette. *Hammering Away - A User's Guide to Sledgehammer for Isabelle/HOL*. Dec. 2021.

[19] C. Barrett, C. L. Conway, M. Deters, L. Hadarean, D. Jovanović, T. King, A. Reynolds, and C. Tinelli. "CVC4." In: *Computer Aided Verification*. Springer Berlin Heidelberg, 2011, pp. 171–177. DOI: `10.1007/978-3-642-22110-1_14`.

[20] *Archive of Formal Proofs*. `https://www.isa-afp.org/`. Accessed July 18, 2022.

[21] P. Lammich and R. Meis. "A Separation Logic Framework for Imperative HOL." In: *Archive of Formal Proofs* (Nov. 2012). `https://isa-afp.org/entries/Separation_Logic_Imperative_HOL.html,` Formal proof development. ISSN: 2150-914x.

[22] T. Nipkow and M. Wenzel. "The Supplemental Isabelle/HOL Library." In: *Isabelle/HOL sessions/HOL-Library* (). `https://isabelle.in.tum.de/library/HOL/HOL-Library/Sublist.html`.